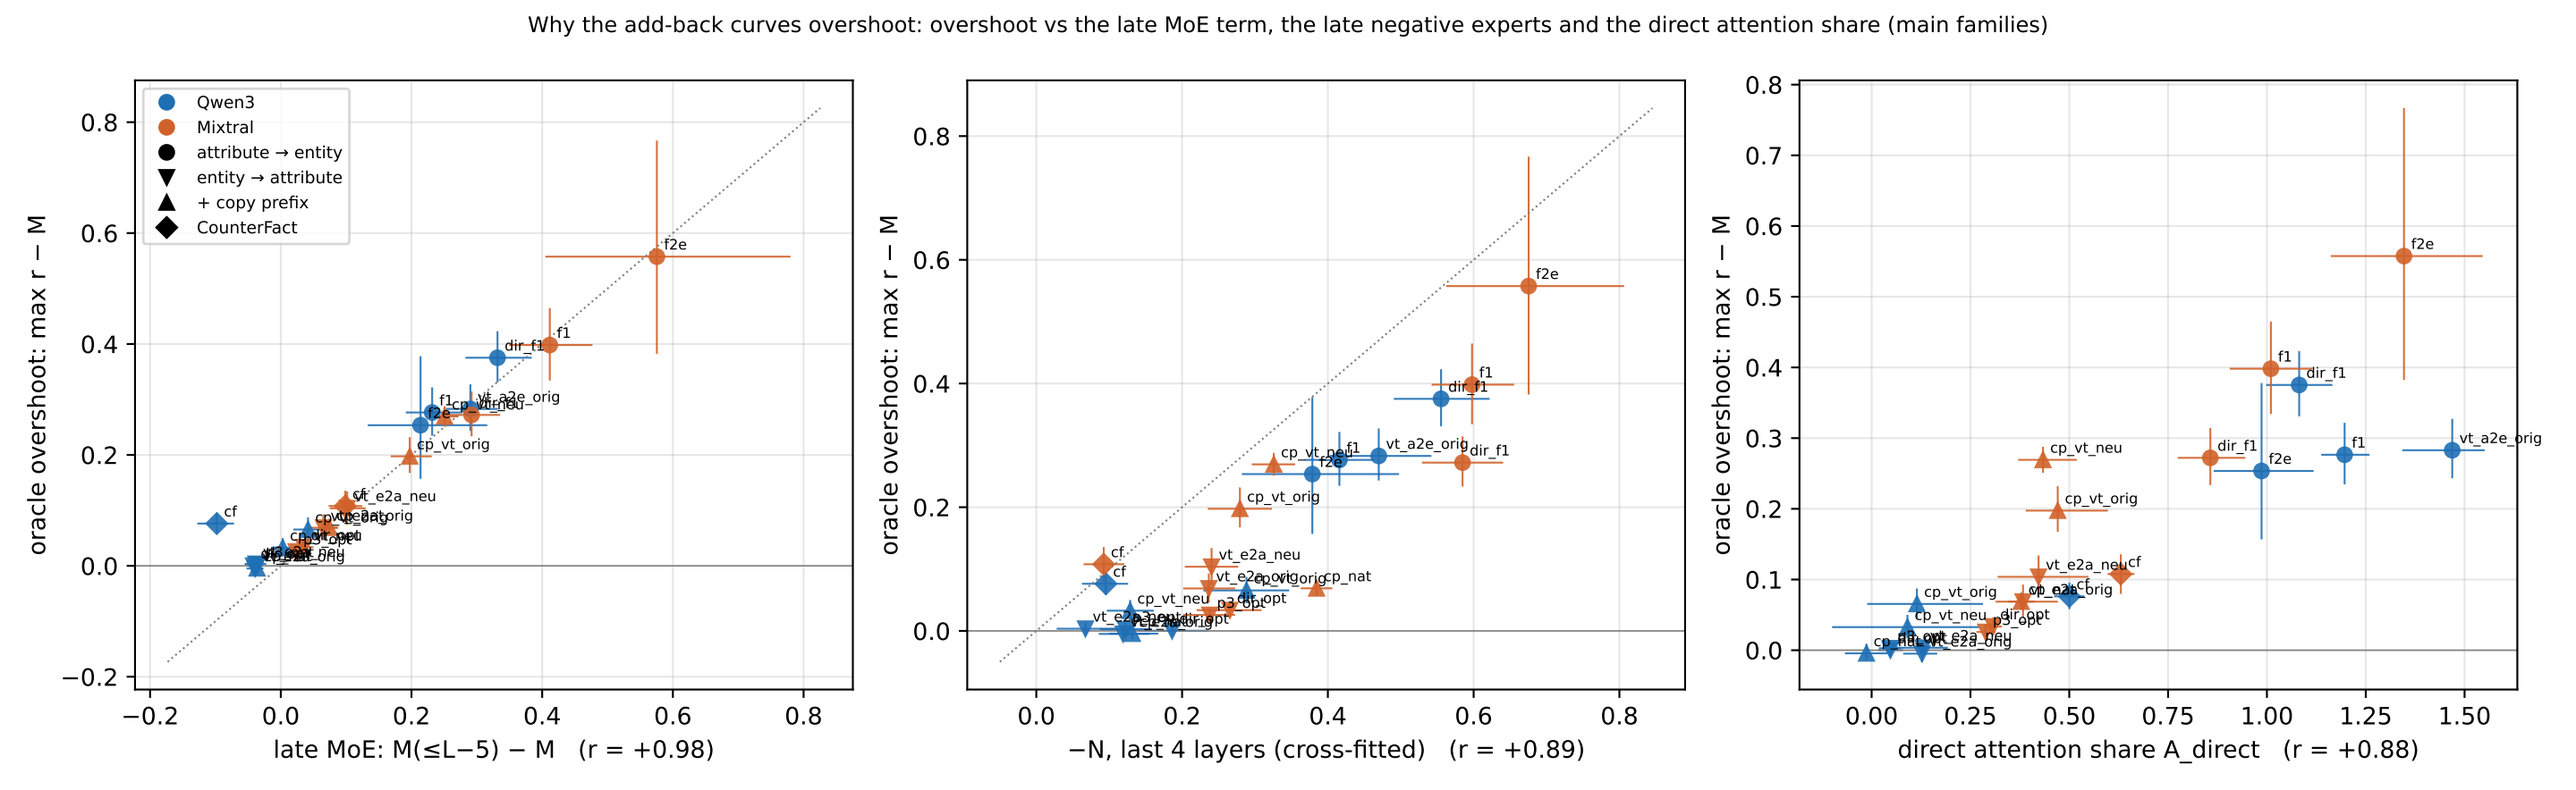
Why the add-back curves overshoot: overshoot vs the late MoE term, the late negative experts and the direct attention share (main families)
Qwen3
Mixtral
attribute → entity
entity → attribute
+ copy prefix
CounterFact
oracle overshoot: max r − M
late MoE: M(≤L−5) − M (r = +0.98)
−N, last 4 layers (cross-fitted) (r = +0.89)
direct attention share A_direct (r = +0.88)
f2e
f1
dir_f1
cp_vt_orig
vt_a2e_orig
vt_e2a_neu
cp_vt_neu
cf
cp_nat
p3_opt
dir_opt